

Write Your Own DNS Resolver

A CS 324 Lab

1 Overview

This lab will provide you a basic understanding of socket programming and will help you understand the complexities and importance of Internet protocols and agreement.

In this lab, you will be building a basic *DNS stub resolver*. A stub resolver is the library on your computer that an application, such as your Web browser or email client, uses to translate a domain name, e.g., `www.example.com` to an IP address, e.g., `192.0.2.1`.

The DNS resolver is considered a *stub* because the extent of its functionality is to:

- format the question into a DNS query of the proper format;
- send the query to a designated *recursive* DNS server;
- wait for and receive the response from the server; and
- extract the answer from the response and return it to the caller.

The recursive DNS server does the hard work of tracking down the answer by issuing its own queries to (sometimes many) other DNS servers.

In this lab the effort will be in communicating effectively with an actual DNS server by implementing the protocol.

2 Downloading the assignment

The files for the lab will be made available from a link in the “assignment” page for the lab on the course site on Learning Suite, as a single archive file, `dnsresolverlab-handout.tar`.

Start by copying `dnsresolverlab-handout.tar` to a protected Linux directory in which you plan to do your work. Then give the command

```
linux> tar xvf dnsresolverlab-handout.tar
```

This will create a directory called `dnsresolverlab-handout` that contains a number of files. You will be modifying the file `resolver.c`.

You can build the executable program by running the following:

```
linux> make clean
linux> make
```

The program is run by supplying the domain name and server as command-line arguments, like this:

```
linux> ./resolver www.example.com 8.8.8.8
93.184.216.34
```

(Note that `8.8.8.8` is a public DNS recursive resolver. Another server to test against is the BYU CS DNS resolver, `128.187.80.20`.)

Its output should match that of the reference implementation, `resolver-ref`:

```
linux> ./resolver-ref www.example.com 8.8.8.8
93.184.216.34
```

3 Description

There are various tasks associated with building a DNS stub resolver.

3.1 Building a DNS Query Message

The first task is organizing the various components of the DNS query into a *DNS query message*. The query message is simply a sequence of bytes that are transmitted to the server over a UDP socket. In C, you can build this sequence of bytes by populating an array of type `unsigned char`. Building the contents of that array is a matter of organizing and formatting the query components according to the DNS protocol specification. There are two major parts to the DNS query message: the DNS header; and the DNS question section, which contains the name and type being resolved).

The DNS header and query are organized as shown in the following page:

<http://www.networksorcery.com/enp/protocol/dns.htm>

in the sections under “DNS header:” and “Query. Variable length”, respectively. Note that the diagrams show values in rows of 32 *bits* (i.e., four bytes). Thus, the entire header for a DNS query message will take up 12 bytes (i.e., three rows).

For example, a query message for `www.example.com` looks like this:

```
27 d6 01 00
```

```
00 01 00 00
00 00 00 00
03 77 77 77
07 65 78 61
6d 70 6c 65
03 63 6f 6d
00 00 01 00
01
```

This is equivalent to:

```
unsigned char msg[] = {
0x27, 0xd6, 0x01, 0x00,
0x00, 0x01, 0x00, 0x00,
0x00, 0x00, 0x00, 0x00,
0x03, 0x77, 0x77, 0x77,
0x07, 0x65, 0x78, 0x61,
0x6d, 0x70, 0x6c, 0x65,
0x03, 0x63, 0x6f, 0x6d,
0x00, 0x00, 0x01, 0x00,
0x01
};
```

In the above example, the following are the header values (beginning at offset 0):

- Identification (query ID) (16 bits): 0x27d6 This value is arbitrary and is typically randomly-generated when the query is formed (note that network protocols use *network byte order*, which is equivalent to Big-endian (most significant bytes are on the left). You can use the `srandom()` and `random()` functions to generate a random integer to fill these bytes (see `man random` for more).
- QR flag (1 bit): 0 (query)
- Opcode (4 bits): 0 (standard query)
- Flags (7 bits): The RD bit is set (1); all others are cleared (0).
- Return Code (4 bits): 0 (NOERROR)
- Total questions (16 bits): 1 (again, note the network byte ordering, most significant bytes on the left)
- Total answer records (16 bits): 0 (this is a query, not an answer, so there are no records in the answer section)
- Total authority records (16 bits): 0 (same)
- Total additional records (16 bits): 0

And the question section contains the query (beginning at byte offset 12):

- Query name: `www.example.com`, encoded as follows (encoding method shown later):

```
03 77 77 77
07 65 78 61
6d 70 6c 65
03 63 6f 6d
00
```

- Type: 1 (type A for address)
- Class: 1 (IN for Internet class)

Note that in this lab, you will *only* be issuing queries for type 1 (A) and class 1 (IN). The header will largely look the same for all queries produced by your resolver—with only the Query ID changing (to random values). Similarly, the question section will look the same, with only the query name changing.

The `create_dns_query()` function has been declared in `resolver.c` and is intended as a helper function to build the `unsigned char` array which will be transmitted to the server. Other helper functions include `name_ascii_to_wire()` and `rr_to_wire()`, also declared in `resolver.c` but left as an exercise to you to define. Note also the `typedef` declarations for `dns_rr` and `struct dns_answer_entry` and the various DNS fields (e.g., `dns_rr_type`). You are welcome to use these functions and types if they are helpful to you, but you are not required to.

The following functions might be useful for your message composition: `strtok()`, `strlen()`, `strncpy()`, `memcpy()`, and (in some cases) `malloc`. Please remember that the string-related functions (e.g., `strlen()`) should be used *only* on null-terminated strings.

3.2 Sending a DNS message and Receiving a Response

Once the bytes in an array of `unsigned char` have been formatted as a DNS query message, their content can be transmitted “on the wire” to a DNS recursive resolver to respond to your query. The `send_recv_message()` helper function is declared to help with sending the message and receiving the response. Note that for the purposes of this lab, the address family will be `AF_INET` (i.e., IPv4), the protocol will be UDP (type `SOCK_DGRAM`), and the destination port is 53 (the standard, well-known DNS port), but an alternate port is optionally passed in from the command line.

Useful structures include the `struct sockaddr_in` data structure (see `man inet_pton`), `socket()` (to create a socket), `connect()` (to indicate which destination IP address and port datagrams should be sent—with UDP there really is no “connection”), `send()` (to send data to a socket), `recv()` (to receive data from a socket), `htons()` (to convert a short to network order), `inet_addr(server)` (for converting a string IP address to bytes). See the `man` pages for each of these functions.

Also note, that because the socket you will use is of type `SOCK_DGRAM`, you will send your query with only a *single* call to `send()` and receive your response with only a *single* call to `recv()`. This is because each `send()` or `recv()` sends or receives at most one datagram.

3.3 Extract an Answer from a DNS Response

Having received a DNS message response from the server, stored in an array of `unsigned char`, the next step is to extract and decode the useful information (i.e., define the `get_answer_address()` function). For the purposes of this lab, the useful parts are the answer count and the resource records in the answer section. You will want to extract each resource record from the section and determine whether or not it is the one you are looking for. Specifically, you will implement the following algorithm (provided in pseudo code):

```
set qname to the initial name queried
(i.e., the query name in the question section)
for each resource record (RR) in the answer section:
    if the owner name of RR matches qname and the type matches the qtype:
        extract the address from the RR, convert it to a string, and add it
        to the result list
    else if the owner name of RR matches qname and the type is (5) CNAME:
        the name is an alias; extract the canonical name from the RR rdata,
        and set qname to that value, and add it to the result list
return NULL (no match was found)
```

The resource records in the answer section are different than the Query in the Question Section, in that they include a time-to-live (TTL) value, Rdata (e.g., the IP address for resource records of type A), and Rdata length:

<http://www.networksorcery.com/enp/protocol/dns.htm>

See the section under “Resource Record. Variable length”.

A response to a query for `www.example.com` might look like this:

```
27 d6 81 80
00 01 00 01
00 00 00 00
03 77 77 77
07 65 78 61
6d 70 6c 65
03 63 6f 6d
00 00 01 00
01 c0 0c 00
01 00 01 00
01 01 82 00
04 5d b8 d8
22
```

This is equivalent to:

```

unsigned char msg[] = {
0x27, 0xd6, 0x81, 0x80,
0x00, 0x01, 0x00, 0x01,
0x00, 0x00, 0x00, 0x00,
0x03, 0x77, 0x77, 0x77,
0x07, 0x65, 0x78, 0x61,
0x6d, 0x70, 0x6c, 0x65,
0x03, 0x63, 0x6f, 0x6d,
0x00, 0x00, 0x01, 0x00,
0x01, 0xc0, 0x0c, 0x00,
0x01, 0x00, 0x01, 0x00,
0x01, 0x01, 0x82, 0x00,
0x04, 0x5d, 0xb8, 0xd8,
0x22
};

```

In the above example, the following are the header values:

- Identification (query ID) (16 bits): 0x27d6 Matches the query ID of the query.
- QR flag (1 bit): 1 (response)
- Opcode (4 bits): 0 (standard query)
- Flags (7 bits): The RD and RA bits are set (1); all others are cleared (0).
- Return Code (4 bits): 0 (NOERROR)
- Total questions (16 bits): 1
- Total answer records (16 bits): 1
- Total authority records (16 bits): 0
- Total additional records (16 bits): 0

The question section is identical to that of the DNS query (see above). The answer section contains a single resource record with the following values:

- Owner name (variable): `www.example.com`, encoded using compression encoding (encoding method shown later):

```
c0 0c
```

- Type (16 bits): 1 (type A for address)
- Class (16 bits): 1 (IN for Internet class)

- TTL (32 bits): 0x00010182 or 65922 (about 18 hours)
- Rdata length (16 bits): 4 (an IPv4 address is four bytes)
- Rdata (variable—specified by the Rdata length field): 0x5db8d822 or 93.184.216.34. These are bytes comprising the IP address corresponding to the owner name.

The `name_ascii_from_wire()` and `rr_from_wire()` helper functions can be defined by you to help with this process. Additionally, the externally-defined functions `malloc()`, `memcpy()`, `inet_ntop()`, and `strcmp()` might be useful. Please remember that the string-related functions (e.g., `strcmp()`) should *only* be used on null-terminated strings.

3.4 Name Encoding and Decoding

Encoding a domain name for a DNS message is explained here:

http://www.tcpiipguide.com/free/t_DNSNameNotationandMessageCompressionTechnique.htm

The DNS query above has an example of this, encoding `www.example.com` thus:

```
03 77 77 77 07 65 78 61 6d 70 6c 65 03 63 6f 6d 00
  w  w  w      e  x  a  m  p  l  e      c  o  m
```

The length (in bytes) of each label precedes the characters comprising the label itself (e.g., the value 3 precedes the three bytes representing the ASCII values `www`). A 0 in the length byte always indicates the end of the labels.

Decoding is similar, but there is a special case called *compression*, which is explained here:

http://www.tcpiipguide.com/free/t_DNSNameNotationandMessageCompressionTechnique-2.htm
http://www.tcpiipguide.com/free/t_DNSNameNotationandMessageCompressionTechnique-3.htm

Note that the DNS response above uses compression for `www.example.com`:

```
c0 0c
```

When the two most significant bits (bits 6 and 7) are set in the first byte—and thus the value is ≥ 192 (i.e., $\geq 0xc0$)—it is an indicator that the byte doesn't represent length; rather it indicates that the *next* byte is a pointer (the offset in the DNS message) to where the sequence of next labels for the name are found. That is, the subsequent byte (after the offset) is no longer part of the name. This pattern is followed until the length byte for a label is 0, indicating the end of the name. In the above case, the byte values are `0xc0` and `0x0c`. Thus, the first indicates that the second is a pointer ($0xc0 = 192_d \geq 192_d$), and the byte value of the second, `0x0c` (12 in decimal), is the offset at start of the question section. Indeed the owner name of the resource record in the answer section is the same as the query name!

3.5 Other Notes

Note that you don't have to use the declared helper functions. The only required function to be defined as declared is `dns_answer_entry *resolve()`. But we hope that the declarations and descriptions in the source will help you on your way.

There are two functions, `print_bytes()` and `canonicalize_name()` that are provided for your benefit to help with development and troubleshooting.

The encoding of the *root* name, represented in string form as ".", is a single byte, `0x00` (note that this is the final byte in encodings of all other domain names because the root is the ancestor of all domain names).

In the case where a DNS name is an alias for another name, a record of type 5 (CNAME) exists. You must include the CNAME target (i.e., the name in the Rdata of the CNAME record) in the list of names returned. It is encoded just as the owner name of the resource record is.

Sometimes there are multiple IP addresses returned for a given name (see `byu.edu`, for example). In this case, all entries must be returned.

You might find it useful to use the command-line tool, `dig`, to issue queries against servers and see a more textual representation of the responses. For example:

```
linux> dig @8.8.8.8 www.example.com
```

4 Testing

You are invited to test your code against multiple DNS resolvers, including `8.8.8.8` and `128.187.80.20` (from the CS network). The following domain names can be used to test:

- `byu.edu` (simple response)
- `www.byu.edu` (single CNAME record, single domain)
- `i-dont-exist.byu.edu` (name that doesn't exist—no answer records)
- `www.intel.com` (multiple CNAME records in response, multiple domains chased for compression)
- `.` the root domain name (note that it won't have any A records, but it shouldn't bomb out)

Also, we have provided in the handout the autograder that we will use to grade your resolver. The autograder will create a proxy DNS server that will handle all DNS requests for grading.

To use the autograder, you can run:

```
linux> make test
```

Note: the autograder requires python 3 and the `dnspython` library. It has been tested on the CS lab machines, but if you are developing on other systems, you might need to install these for it to work properly.

The autograder invokes your resolver to resolve the following domain names:

- `byu.edu`
- `www.byu.edu`
- `.` (root domain)
- `casey.byu.edu`
- `www.abcdefghijklmnopqrstuvwxyzabcdefghijklmnopqrstuvwxyzabcdefghijklmnopqrstuvwxyz`
- `sandia.gov`
- `www.sandia.gov`
- `foobarbaz-not-exist.byu.edu`
- `www.intel.com`

The autograder also checks for memory leaks and compiler warnings. In order to get full credit for the lab, your code should compile with no warnings, and there should be no memory leaks (i.e., you should be calling `free()` on any memory that you `malloc()`).

5 Grading

The following is the point breakdown:

- 25 points for a well-formed DNS query message
- 20 points for successfully sending the query and receiving the response
- 25 points for successfully finding the answer in the answer section
- 10 points for handling CNAME records properly
- 10 points for handling names that don't resolve
- 5 points for handling the root name properly
- 3 compiles without any warnings
- 2 no memory leaks

The maximum number of points is 100.

6 Handing in Your Work

To submit your lab, please upload your `resolver.c` file to the assignment page corresponding to the DNS resolver lab on Learning Suite.